Q: Did you participate in any discussions about user interactions with the model?

A: I joined the team in the middle of the project, so I did not participate much in the initial discussions. I noticed that client representatives could actively suggest model parameter adjustments. Another discussed topic was to keep information about the model's operation private from end users. We did this to prevent them from learning how to manipulate the model in their favor.

Q: When you joined the project, were the definitions related to the model's objectives and the system architecture passed on to you?

A: No, I became aware of them during the project. I would ask questions to the data science team when I had doubts. There was no formal passage of knowledge but instead explanations on demand.

Q: Do you perceive that as a difficulty?

A: Yes. When we met with client representatives and data scientists to map the data required by the model, I was unsure if the data we requested was correct, since I did not know what the data scientists expected for the model input.

In these meetings, I noticed a mismatch regarding the participants' understanding of the data. For example, I was expecting them to be in a particular format. Yet, the data scientists were expecting them in another form, and what the client representatives understood differed from what the data scientists were expecting. This situation provoked changes throughout the development of the API used to consume the model.

Q: Did you participate in discussions about obtaining the data for the model?

A: I was not involved in this discussion. As far as I know, data were not obtained through the cloud or online access. The data scientists received files from the client representatives for training and validating the model. Still, the data were placed in a secure area since it was sensitive.

Q: Did you participate in any data-related activity?

A: No, everything was in the hands of the data science team.

Q: Did you participate in the definition of how the system would consume the model?

A: I developed the services that use the model consumption API. I do not know if it is easy to change the model.

Q: How was the discussion about the storage of the model artifacts?

A: I did not participate in that conversation. I think all artifacts are in a Git repository.

Q: Was any data scientist involved in the publication of the model?

A: No, we are responsible for deploying the model consumption API. The deployment of this service, as well as the other system services, is automated through a CI/CD pipeline.

Q: Was there any discussion about model updates and implementing incremental learning?

A: We discussed this, it was a feature raised by client representatives, but the discussion did not go very far. For now, we initially defined that retraining the model with new data would be a manual task. The data scientists also participated in these discussions.

Q: How was the discussion regarding model monitoring?

A: I know the data science team will send model results to client representatives to assess if they are as expected, which is also done manually. We are not in charge of that. However, if there is a problem with the infrastructure of any service, the responsibility falls to our team.

Q: Were security and privacy issues considered when building the system?

A: We tried to protect communication with external APIs and services as much as possible, using JWT tokens that expire after some time. Private data consumed by internal services are persisted in databases and are not exposed to other users.

Q: How was the relationship with the data scientists during the integration of the model with the rest of the system?

A: We were very separated, and I did not like that. We did not know much about the model. It was like a "black box" that we did not get involved much. That's how it happened, and I do not know if it was supposed to be like that. Even with a well-defined API, things that were obvious to the data science team were unclear to us. We did not participate in creating the model. We only developed the services that consumed it, so we did not know what was being done. This proved to be a problem when we met with the client to define the data we needed from their external APIs.

This separation got in the way of integrating the system with the rest of the model. I did not have the necessary knowledge to analyze if the data was correct and what fields were required or optional. Problems only appeared when we started testing and integrating the system with the client's external APIs. Only then we noticed data was either missing or in the wrong format. If

the teams were not so distant, we could have anticipated these problems.

Q: Is the integration of the model with the rest of the system well-documented?

A: No, it is not documented well enough. We currently have the model's output and input data documented. But, for example, in the middle of this integration, there is a mapper that converts data to the format expected by the model. We could have documented this conversion better.